

```
> #2(a)
> #Forward selection
> data(seatpos, package="faraway")
> lmod=lm(Seated~1, data=seatpos)
> indep.vars=~Age+Weight+Ht+hipcenter+Leg
> add1(lmod, indep.vars, test="F")
Single term additions
```

Model:

Seated ~ 1

	Df	Sum of Sq	RSS	AIC	F value	Pr(>F)
<none>			899.93	122.260		
Age	1	26.07	873.86	123.143	1.074	0.307
Weight	1	541.40	358.54	89.289	54.361	1.071e-08 ***
Ht	1	775.39	124.54	49.109	224.129	< 2.2e-16 ***
hipcenter	1	481.22	418.71	95.185	41.375	1.845e-07 ***
Leg	1	593.24	306.69	83.354	69.635	6.172e-10 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
> lmod=update(lmod, .~. +Ht)
> add1(lmod, indep.vars, test="F")
Single term additions
```

Model:

Seated ~ Ht

	Df	Sum of Sq	RSS	AIC	F value	Pr(>F)
<none>			124.54	49.109		
Age	1	6.7957	117.75	48.977	2.0200	0.1641
Weight	1	0.1238	124.42	51.071	0.0348	0.8531
hipcenter	1	0.2657	124.28	51.028	0.0748	0.7861
Leg	1	5.5300	119.01	49.383	1.6263	0.2106

```
> #Ht remains in the final model
```

```
> #2(b)
> #Using backward selection
> data(seatpos, package="faraway")
> lmod=lm(Seated~Age+Weight+Ht+hipcenter+Leg, data=seatpos)
> indep.vars=~Age+Weight+Ht+hipcenter+Leg
> drop1(lmod, test="F")
```

Single term deletions

Model:

```
Seated ~ Age + Weight + Ht + hipcenter + Leg
      Df Sum of Sq  RSS   AIC F value    Pr(>F)
<none>                 111.16 52.789
Age      1      7.396 118.56 53.236  2.1291    0.1543
Weight   1      1.804 112.96 51.400  0.5192    0.4764
Ht        1     114.451 225.61 77.687 32.9473 2.307e-06 ***
hipcenter 1      0.278 111.44 50.884  0.0800    0.7791
Leg       1      4.073 115.23 52.156  1.1726    0.2870
```

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
> lmod=update(lmod, ~. -hipcenter)
> drop1(lmod, test="F")
```

Single term deletions

Model:

```
Seated ~ Age + Weight + Ht + Leg
      Df Sum of Sq  RSS   AIC F value    Pr(>F)
<none>                 111.44 50.884
Age      1      7.133 118.57 51.241  2.1122    0.1556
Weight   1      1.853 113.29 49.510  0.5488    0.4640
Ht        1     120.960 232.40 76.813 35.8198 1.008e-06 ***
Leg       1      5.051 116.49 50.568  1.4957    0.2300
```

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
> lmod=update(lmod, ~. -Weight)
> drop1(lmod, test="F")
```

Single term deletions

Model:

```
Seated ~ Age + Ht + Leg
      Df Sum of Sq  RSS   AIC F value    Pr(>F)
<none>                 113.29 49.510
```

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
> lmod=update(lmod, ~. -Weight)
```

```
> drop1(lmod, test="F")
```

Single term deletions

Model:

Seated ~ Age + Ht + Leg

	Df	Sum of Sq	RSS	AIC	F value	Pr(>F)
<none>			113.29	49.510		
Age	1	5.723	119.01	49.383	1.7176	0.1988
Ht	1	176.768	290.06	83.235	53.0499	1.96e-08 ***
Leg	1	4.457	117.75	48.977	1.3377	0.2555

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
> lmod=update(lmod, ~. -Leg)
```

```
> drop1(lmod, test="F")
```

Single term deletions

Model:

Seated ~ Age + Ht

	Df	Sum of Sq	RSS	AIC	F value	Pr(>F)
<none>			117.75	48.977		
Age	1	6.80	124.54	49.109	2.02	0.1641
Ht	1	756.12	873.86	123.143	224.75	<2e-16 ***

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
> lmod=update(lmod, ~. -Age)
```

```
> drop1(lmod, test="F")
```

Single term deletions

Model:

Seated ~ Ht

	Df	Sum of Sq	RSS	AIC	F value	Pr(>F)
<none>			124.54	49.109		
Ht	1	775.39	899.93	122.260	224.13	< 2.2e-16 ***

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
> #Ht is the only remaining variable
```

1

```

> #2(c)
> data(seatpos, package="faraway")
> x=model.matrix(Seated~Age+Weight+Ht+hipcenter+Leg-1, data=seatpos)
> y=seatpos$Seated
> bestmods=leaps(x,y,nbest=1)
> bestmods
$which
      1      2      3      4      5
1 FALSE FALSE TRUE  FALSE FALSE
2  TRUE FALSE TRUE  FALSE FALSE
3  TRUE FALSE TRUE  FALSE  TRUE
4  TRUE  TRUE TRUE  FALSE  TRUE
5  TRUE  TRUE TRUE   TRUE   TRUE


$label
[1] "(Intercept)" "1"      "2"      "3"      "4"
[6] "5"

$size
[1] 2 3 4 5 6

$Cp
[1] 1.852984 1.896690 2.613524 4.079995 6.000000

> Cpplot(bestmods)
> #From the graph, we can see the model with the best number of parameters is two, it only
contains the variable weight, since it's closer to the line p and below it. The one with 6
parameter is also possible but we don't choose it because it has too many variables.
>
>
>
>
>
>
>
>
>

```



```
> #2(d)
> data(seatpos, package="faraway")
> X=as.matrix(seatpos[,c(1,2,4,8,9)])
> Y=seatpos[,5]
> mod_cv=cv.glmnet(X,Y)
> model.lasso=glmnet(X,Y,lambda=mod_cv$lambda.min)
> blasso=model.lasso$beta
> blasso
5 x 1 sparse Matrix of class "dgCMatrix"
           s0
Age      -0.022086496
Weight    0.005602745
Ht        0.380163808
Leg        .
hipcenter  .
> #We pick Age, Weight and Ht in the Lasso model.
```

```

> #3(a)
> beta=c(1,-1,1,-1,0,1,-1,1,-1,0,1,-1,1,-1,0,1,-1,1,-1,0)
> X=rmvnorm(100,mean=rep(0,25),sigma=diag(0.7,25,25)+matrix(0.3,25,25))
> #3(b)
> errors=rnorm(100,0,1)
> Y=X%%beta+errors
> #3(c)
> #Ridge
> cv.ridge=cv.glmnet(X,Y,alpha=0)
> model.ridge=glmnet(X,Y,lambda=cv.ridge$lambda.min, alpha=0)
> bridge=as.matrix(model.ridge$beta)
> rmse=mse(beta,bridge)
> #Lasso
> mod_cv=cv.glmnet(X,Y)
> model.lasso=glmnet(X,Y,lambda=mod_cv$lambda.min)
> blasso=as.matrix(model.lasso$beta)
> lmse=mse(beta,blasso)
> #Least Squares
> beta_i=solve((t(X))%%X)%%(t(X))%%Y
> lsmse=mse(beta_i, beta)
> rbind(rmse,lmse,lsmse)
      [,1]
rmse 0.02330827
lmse 0.01319720
lsmse 0.01315594
> #As shown by the results, the mse for the beta values for each estimators are similar.

```

```

> #Repeat ridge 100 times and take the average
> mse_100times_ridge=replicate(100,{cv.ridge=cv.glmnet(X,Y,alpha=0)
+ model.ridge=glmnet(X,Y,lambda=cv.ridge$lambda.min, alpha=0)
+ bridge=as.matrix(model.ridge$beta)
+ rmse=mse(beta,bridge)})
> rmse_100=mean(mse_100times_ridge)
> #Repeat lasso 100 times and take the average
> mse_100times_lasso=replicate(100,{mod_cv=cv.glmnet(X,Y)
+ model.lasso=glmnet(X,Y,lambda=mod_cv$lambda.min)
+ blasso=as.matrix(model.lasso$beta)
+ lmse=mse(beta,blasso)})
> lmse_100=mean(mse_100times_lasso)
> #Repeat least squares 100 times and take the average
> mse_100times_least_squares=replicate(100, {beta=c(1,-1,1,-1,0,1,-1,1,-1,0,1,-1,1,-1,0,1,-
1,1,-1,0,1,-1,1,-1,0)
+ X=rmvnorm(100,mean=rep(0,25),sigma=diag(0.7,25,25)+matrix(0.3,25,25))
+ errors=rnorm(100,0,1)
+ Y=X%%beta+errors
+ beta_i=solve((t(X))%%X)%%(t(X))%%Y
+ mse(beta_i, beta)})
> lsmse_100=mean(mse_100times_least_squares)
> rbind(rmse_100,lmse_100,lsmse_100)
      [,1]
rmse_100 0.02342235
lmse_100 0.01377160
lsmse_100 0.01852426
> #Again, the results are similar.

```

```

> #Again, the results are similar, but overall lasso has a smaller mse so it's a slightly better model.
> #3(c)-repeat
> beta=c(1,-1,1,-1,1,-1,1,-1,1,-1,1,-1,1,-1,1,-1,0,0,0,0,0)
> X=rmvnorm(100,mean=rep(0,25),sigma=diag(0.7,25,25)+matrix(0.3,25,25))
> errors=rnorm(100,0,1)
> Y=X%%beta+errors
> #Ridge
> cv.ridge=cv.glmnet(X,Y,alpha=0)
> model.ridge=glmnet(X,Y,lambda=cv.ridge$lambda.min, alpha=0)
> bridge=as.matrix(model.ridge$beta)
> rmse=mse(beta, bridge)
> #Lasso
> mod_cv=cv.glmnet(X,Y)
> model.lasso=glmnet(X,Y,lambda=mod_cv$lambda.min)
> blasso=as.matrix(model.lasso$beta)
> lmse=mse(beta, blasso)
> #Least squares
> beta_i=solve((t(X))%%X)%%(t(X))%%Y
> lsmse=mse(beta_i, beta)
> rbind(rmse,lmse,lsmse)
      [,1]
rmse 0.02475615
lmse 0.02015487
lsmse 0.01763795
> #All estimators have similar results of Mean Squared Errors.

```



```

> #Repeat Ridge 100 times and take the average
> mse_100times_ridge=replicate(100,{cv.ridge=cv.glmnet(X,Y,alpha=0)
+ model.ridge=glmnet(X,Y,lambda=cv.ridge$lambda.min, alpha=0)
+ bridge=as.matrix(model.ridge$beta)
+ mse(beta, bridge)})
> r_100mse=mean(mse_100times_ridge)
> #Repeat Lasso 100 times and take the average
> mse_100times_lasso=replicate(100,{mod_cv=cv.glmnet(X,Y)
+ model.lasso=glmnet(X,Y,lambda=mod_cv$lambda.min)
+ blasso=as.matrix(model.lasso$beta)
+ mse(beta, blasso)})
> l_100mse=mean(mse_100times_lasso)
> #Repeat Least Squares 100 times and take the average
> mse_100times_least_squares=replicate(100, {beta=c(1,-1,1,-1,1,-1,1,-1,1,-1,1,-1,1,-1,1,-1,1,-1,1,-1,0,0,0,0,0)
+ X=rmvnorm(100,mean=rep(0,25),sigma=diag(0.7,25,25)+matrix(0.3,25,25))
+ errors=rnorm(100,0,1)
+ Y=X%%beta+errors
+ beta_i=solve((t(X))%%X)%%(t(X))%%Y
+ mse(beta_i, beta)})
> ls_100mse=mean(mse_100times_least_squares)
> rbind(r_100mse,l_100mse,ls_100mse)
      [,1]
r_100mse 0.02479899
l_100mse 0.02003675
ls_100mse 0.01892135
> #Results are similar

```